

Week 3 Entropy Project

Sandesh Sadalge

Fri, Sept 12, 2014

Dataset

I put the data into my github directory to make things easier. Need to load a package before I can get it from github with this R-Markdown.

```
## Loading required package: bitops

fileLoc <-
"https://raw.githubusercontent.com/spenglerss/MSDAFall2014/master/IS%20607%20Data%20Acquisition/P1_Entropy/entropy-test-file.csv"
dataset <- read.csv(text = getURL(fileLoc))
```

Obligatory summary information of the dataset. I think str() gives you a little more informative output for our purposes:

```
head(dataset, n=5)

##   attr1 attr2 attr3 answer
## 1     0     A     2      0
## 2     0     A     2      0
## 3     0     A     2      0
## 4     0     A     1      0
## 5     0     A     1      0

summary(dataset)

##      attr1      attr2      attr3      answer
## Min.   :0.0    A:360   Min.   :0.000   Min.   :0.000
## 1st Qu.:0.0    B:340   1st Qu.:0.000   1st Qu.:0.000
## Median :0.0    C:300   Median :1.000   Median :0.000
## Mean   :0.4                Mean   :0.991   Mean   :0.424
## 3rd Qu.:1.0                3rd Qu.:2.000   3rd Qu.:1.000
## Max.   :1.0                Max.   :2.000   Max.   :1.000

str(dataset)

## 'data.frame': 1000 obs. of 4 variables:
## $ attr1 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ attr2 : Factor w/ 3 levels "A","B","C": 1 1 1 1 1 1 1 1 1 1 ...
## $ attr3 : int 2 2 2 1 1 0 1 0 2 0 ...
## $ answer: int 0 0 0 0 0 0 1 0 0 0 ...
```

A) Create the **entropy()** function:

Here is my final function code. Explanations follow.

```
entropy <- function (d)
{
  partitions <- table(d)                                # (1)
```

```

partitions <- partitions / sum(partitions)    # (2)
partitions <- partitions * log2(partitions)   # (3)
entropy <- -sum(partitions)                  # (4)
}

```

(1) Determining categorical partitions and their frequencies:

Initially, the hard part for me was coming up with a way to determine a given vector's categorical partitions and their frequencies.

My first thoughts started down the path of either trying to convert to factor or using *unique()* to get the distinct categorical values and perhaps then using a loop to figure out the frequency. I realized this was probably a very 'un-R' way to do it.

Fortunately, I came across **table()** which seemed to do exactly what I was looking for. The code: `partitions <- table(d)` gives you a table with the distinct values & respective frequencies of the input vector, d. As an example, in our dataset, `table(dataset$answer)` returns:

```

##
##    0    1
## 576 424

```

(2) Thanks to the fact that R is a vectorized language, `partitions <- partitions / sum(partitions)` determines the probability of each of the partitions (based on frequency counts).

Again, in our example on `table(dataset$answer)`, when you do this, you get:

```

##
##      0      1
## 0.576 0.424

```

(3) Similarly, `partitions <- partitions * log2(partitions)` gives you the individual calculation within the summation notation of the Entropy equation, $\text{probability} * \log_2(\text{probability})$.

```

##
##      0      1
## -0.4584 -0.5249

```

(4) And lastly, we just take the negative sum of the terms: `entropy <- -sum(partitions)`

```

## [1] 0.9833

```

B) Create the **infogain()** function

Here's the code:

```

infogain <- function (d, a)
{
  if (all(d==a) == TRUE)                # (1)
  {
    infogain <- 0
  } else
  {
    entropy.d <- entropy(d)              # (2)

    x <- table(a, d)                     # (3)

```

```

y <- (x / rowSums(x)) # (4)
y <- y * ifelse(is.infinite(log2(y)),0,log2(y)) # (5)
z <- rowSums(x) / sum(x) # (6)
entropy.a <- sum(z * -1 * rowSums(y)) # (7)

infogain <- entropy.d - entropy.a
}
}

```

(1) if (all(d==a) == TRUE) This is one of those pieces of code that one puts in after you've started testing and you realize there's a special case that needs to be taken care of. I explain this much more fully in note 3 in the next function creation, decide(). Basically, if the d and a vectors are the same, the infogain should be 0. You don't 'gain' anything for the same partition, right?

Note on explanations below: I wish I could put together a video or something to explain how I'm calculating the probabilities and entropies using tables but I'll try by best to explain using my limited R-Markdown abilities & some pictures I made in Excel:

Say you have a vector **d** with partitions **d.i** where **i = 1 ... N** and another vector **a** with partitions **a.j** where **j = 1 ... K**. Then, the infogain, **I(d,a) = E(d) - E(d,a)**.

You can get E(d) using the first function created which is what I did in **(2)** using entropy(d).

As a reminder, you get the following entropy for entropy(dataset\$answer):

```
## [1] 0.9833
```

But how do you get E(d,a)?

I realized later that I could've created another function called entropy(d,a) to do this but I opted to put the code withint the infogain() function instead.

My solution uses *table()* again to create a 2 way table whose column names are *d.i* and row names are *a.j*. the value (*a.j,d.i*) equals the *frequency of the partitions a.j & d.i*

Something like this:

table(d,a)	d ₁	...	d _i	...	d _N
a ₁	freq(a ₁ ,d ₁)		freq(a ₁ ,d _i)		freq(a ₁ ,d _N)
⋮					
a _j	freq(a _j ,d ₁)		freq(a _j ,d _i)		freq(a _j ,d _N)
⋮					
a _K	freq(a _K ,d ₁)		freq(a _K ,d _i)		freq(a _K ,d _N)

This is what line **(3)** does using table(a, d)

To be less abstract, for d <- dataset\$answer and a <- dataset\$attr1 you get the following table:

```
##
##      0    1
## 0 347 253
## 1 229 171
```

(4) In how I've constructed this table, you will notice that the sum of the rows are the total frequency of the partitions a_j :

table(d,a)	d_1	...	d_i	...	d_N	
a_1	freq(a_1, d_1)		freq(a_1, d_i)		freq(a_1, d_N)	<- sum of row $a_1 = \text{freq}(a_1)$
\vdots						
a_j	freq(a_j, d_1)		freq(a_j, d_i)		freq(a_j, d_N)	<- sum of row $a_j = \text{freq}(a_j)$
\vdots						
a_K	freq(a_K, d_1)		freq(a_K, d_i)		freq(a_K, d_N)	<- sum of row $a_K = \text{freq}(a_K)$

I use this fact to figure out the probabilities $P(a_j|d_i)$ (or is the notation $P(d_i|a_j)$?) simply by dividing each element by the sum of the row. Or in my code, $(x / \text{rowSums}(x))$.

Continuing our example using `dataset$answer` and `dataset$attr1`, you get:

```
##
##           0           1
##    0 0.5783 0.4217
##    1 0.5725 0.4275
```

Notice that the sum for each of the rows is 1.

(5) Now that we have the probabilities, we just multiply by the log base 2 of each one respectively as the entropy equation states.

Here you'll notice that I put a special case using `ifelse` for when `log2()` is infinity. One case where this will happen is if `d` and `a` are the same vector.

Using `dataset$answer` and `dataset$attr1` we get:

```
##
##           0           1
##    0 -0.4569 -0.5253
##    1 -0.4607 -0.5241
```

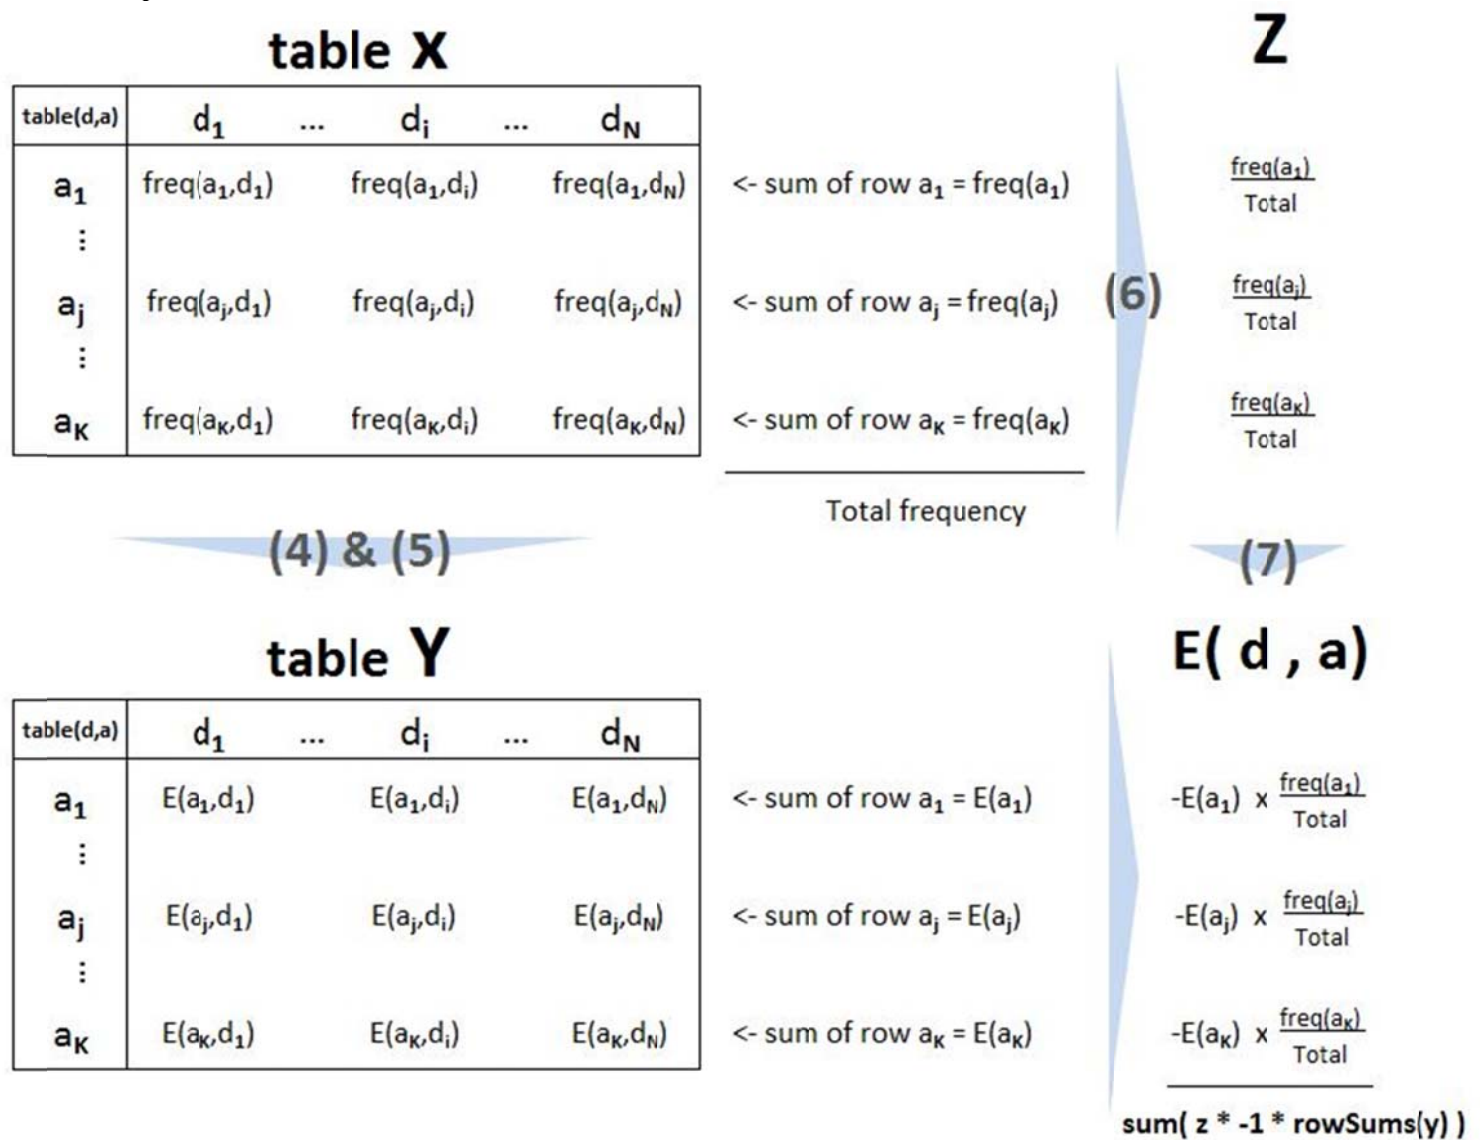
Notice, after line (5), table x has the frequencies and table y has the component pieces of $E(a_j)$!!

Therefore, `rowSums(y)`, will now give you $E(a_i)$. This fact will be useful in **(7)**.

(6) We need one last thing which is the weights for $E(a_j)$. Since table `x` still holds the frequencies, the weights are simply `rowSums(x) / sum(x)` which I put into `z`.

(7)* What does `sum(z * -1 * rowSums(y))` do? Well, since `rowSums(y)` is the entropy for $E(a_j)$ and `z` has the weights for a_j , multiplying the two, taking the negative of it and summing it up will give you the total entropy $E(d,a)$.

The whole process looks like this:



Examples for (6) & (7):

Entropies E(a.j) or rowSums(y)

```
##      0      1
## -0.9822 -0.9848
```

Weights in table z:

```
##      0      1
## 0.6 0.4
```

Above two are multiplied & added and negated to get to the final total for E(d,a):

```
## [1] 0.9832
```

Lastly, the Infogain is E(d) - E(d,a) which is:

```
## [1] 2.412e-05
```

Examples asked for in the assignment (aparently R-Markdown rounds a little differently than the RStudio console):

```
infogain(dataset$answer,dataset$attr1) =  $2.4116 \times 10^{-5}$ 
```

```
infogain(dataset$answer,dataset$attr2) = 0.2599
```

```
infogain(dataset$answer,dataset$attr3) = 0.0024
```

C) Create the **decide()** function

Here's the code:

```
decide <- function(inputDF, col)
{
  x <- apply(inputDF,2,infogain,inputDF[,col])
  decide <- list(max=which.max(x), gains=x)
}
```

After hours of trying to figure out how to do this using *apply()*, ***I finally got it!***

The explanation is pretty easy. You simply *apply()* the function *infogain()* to each of the columns in the inputDF against the reference column. I know the statement is really simple but boy did it take me forever to finally figure out why and how to get it right!

The *apply()* outputs a vector assigned, to x, which allows you to easily lookup the maximum value using *which.max()*.

It's a simple thing then, to return a list with the maximum infogain column number and the list of infogains by column as the assignment asks for.

Oh, and last note:

It was after writing this function that I had to go back to the original *infogain()* function I wrote and put in the special case for checking if the vectors are the same. The reason is that since I iterate through all the columns, I will always try to do the *infogain* for the reference vector and itself. *Infogain* in this case should be 0. (You shouldn't 'gain' anything if you subset back on the same criteria.)

What actually happened in my original *infogain()* function was that *E(d)* gave a valid number but when I tried to do *entropy(d,d)* it rerutned 0. Why? because I put a catch for the infinities associated with $\log_2(0)$ which made $E(d,a) = 0$. Therefore, the $\text{infogain} = E(d) - E(d,a)$ was reverting to simply *E(d)* which was wrong.

After puzzling through it, there were two solutions I could see: Either I do not account for the infinities and allow the *infogain()* function to return NaN or I could put a test in the very begining of the *infogain()* function to check if the vectors are exactly the same. I chose the latter.

Here is the original function I wrote and glad I didn't submit:

```
decide2 <- function(inputDF, col)
{
  decideDF <- data.frame(colnames=colnames(inputDF))
  n <- length(decideDF$colnames)
  infogain.values <- rep(0, n)
```

```

for (i in 1:n)
{
  infogain.values[i] <- infogain(inputDF[,col], inputDF[,i])
}

decideDF <- data.frame(colnames=decideDF$colnames,infogain=infogain.values)

decide2 <- list(max=which.max(infogain.values), gains=decideDF)
}

```

Lastly, let's see if the outputs match the expected ones:

```

(entropy(dataset$answer))
## [1] 0.9833

(infogain(dataset$answer, dataset$attr1))
## [1] 2.412e-05

(infogain(dataset$answer, dataset$attr2))
## [1] 0.2599

(infogain(dataset$answer, dataset$attr3))
## [1] 0.002433

(decide(dataset,4))
## $max
## attr2
##      2
##
## $gains
##      attr1      attr2      attr3      answer
## 2.412e-05 2.599e-01 2.433e-03 0.000e+00

```